On PDE approach to some learning problems

Hong Zhang

Applied Math, Brown University

April 7, 2017

Partial differential equations (PDEs) are used to describe many phenomenons. Many famous PDEs are from physics.

- Heat: heat equation, v tempature

$$\partial_t u = \sum_{i=1}^{3} \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$

Partial differential equations (PDEs) are used to describe many phenomenons. Many famous PDEs are from physics.

- Heat: heat equation, v tempature

$$\partial_t u = \sum_{i=1}^{3} \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$

- Fluid mechanics: Naiver-Stokes equation

$$\partial_t u + (u \cdot \nabla) u - \Delta u = -\nabla p.$$

Partial differential equations (PDEs) are used to describe many phenomenons. Many famous PDEs are from physics.

- Heat: heat equation, v tempature

$$\partial_t u = \sum_{i=1}^{3} \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$

- Fluid mechanics: Naiver-Stokes equation

$$\partial_t u + (u \cdot \nabla)u - \Delta u = -\nabla p.$$

- Quantum mechanics, electrodynamics, image processing,....

Probability and PDEs


Figure: Brownian Motion

Let $g : S_1 \to \mathbb{R}$. Let W be a Brownian motion with $W_0 = x$. W hits $\partial\Omega$ at W_τ. Let $u(x) = \mathbb{E}^x[g(W_\tau)]$. Then

$$\Delta u = 0,$$

and

$$u = g \quad \text{on} \quad S_1.$$

Online Learning: Examples

Figure: Spam Detection

- emails arrive one by one

Online Learning: Examples

Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: spam or not

Online Learning: Examples

Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: spam or not
- sometimes get feedback from users

Online Learning: Examples

Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: spam or not
- sometimes get feedback from users
- make better detection next time

Online Learning: Example 2

Figure: Recommendation

- users visit Netflix /Amazon/...

Online Learning: Example 2

Figure: Recommendation

- users visit Netflix /Amazon/…
- movies/products/…are recommended

Online Learning: Example 2

Figure: Recommendation

- users visit Netflix /Amazon/...
- movies/products/... are recommended
- users click or not

Online Learning: Example 2

Figure: Recommendation

- users visit Netflix /Amazon/...
- movies/products/...are recommended
- users click or not
- make better recommendation next time

Motivation

- Batch learning:
  - distribution fixed over time (training and test)
  - i.i.d. assumptions
  - a fixed predictor

Motivation

- Batch learning:
    - distribution fixed over time (training and test)
    - i.i.d. assumptions
    - a fixed predictor

- Online learning:
    - no distribution assumption
    - suitable for changing or adversarial environment
    - a natural model for many applications, mixing training and test
    - memory efficient

The expert problem

- Need: to make a decision every day
  - which stock to invest in
  - which route to drive home
  - . . .

The expert problem

- Need: to make a decision every day
  - which stock to invest in
  - which route to drive home
  - . . .
- Given: advice from a set of experts
  - financial adviser
  - Route 1, Route I-95, Route 295,. . .
  - . . .

The expert problem

- Need: to make a decision every day
  - which stock to invest in
  - which route to drive home
  - ...
- Given: advice from a set of experts
  - financial adviser
  - Route 1, Route I-95, Route 295,...
  - ...
- Question: <u>How to combine advice and make smart choices?</u>

Gentle Start: a simple example

- Set up:
    - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.

Gentle Start: a simple example

- Set up:
  - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.
  - N experts provide advice $(f_{1,t}, f_{2,t}, \ldots, f_{N,t})$

Gentle Start: a simple example

- Set up:
  - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.
  - N experts provide advice $(f_{1,t}, f_{2,t}, \ldots, f_{N,t})$
  - A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t

Gentle Start: a simple example

- Set up:
    - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.
    - N experts provide advice $(f_{1,t}, f_{2,t}, \ldots, f_{N,t})$
    - A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
    - True y_t is revealed and find if $\hat{p}_t = y_t$.

Gentle Start: a simple example

- Set up:
  - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.
  - N experts provide advice $(f_{1,t}, f_{2,t}, \ldots, f_{N,t})$
  - A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
  - True y_t is revealed and find if $\hat{p}_t = y_t$.
  - Know that one of the experts makes no mistakes.

Gentle Start: a simple example

- Set up:
  - Predict an unknown sequence $y_1, y_2, \ldots$, where $y_i \in \{0, 1\}$.
  - N experts provide advice $(f_{1,t}, f_{2,t}, \ldots, f_{N,t})$
  - A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
  - True y_t is revealed and find if $\hat{p}_t = y_t$.
  - Know that one of the experts makes no mistakes.

- Goal: bound the number of the mistakes made by the forescaster

Halving Algorithm

Algorithm:

- at time t, expert i has weight w_i^t

Halving Algorithm

Algorithm:

- at time t, expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$

Halving Algorithm

Algorithm:

- at time t, expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority

Halving Algorithm

Algorithm:

- at time t, expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority
- update each wrong expert $w_i^t = 0$

Halving Algorithm

Algorithm:

- at time t, expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority
- update each wrong expert $w_i^t = 0$

Conclusion: the number of mistakes made by the forecaster is bounded by $[\log_2 N]$.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.
- When the forecaster make her mth mistakes, at least half of the experts have been always correct make their first mistake

$$w_m \leq \frac{w_{m-1}}{2} \implies w_m \leq \frac{w_0}{2^m}.$$

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.
- When the forecaster make her mth mistakes, at least half of the experts have been always correct make their first mistake

$$w_m \leq \frac{w_{m-1}}{2} \implies w_m \leq \frac{w_0}{2^m}.$$

- $w_i = 1$ for some i

$$m \leq [\log_2 N].$$

More general setting:

- Do not know there is one expert who makes no mistakes.

More general setting:

- Do not know there is one expert who makes no mistakes.
- Weighted Majority Algorithm:
  - at any time, expert i has weight w_i^t.
  - originally, $w_i^t = 1, \forall i \in [1, N]$.
  - prediction according to weighted majority.
  - weight of each wrong expert updated ($\beta \in (0, 1)$)

$$w_i^{t+1} \leftarrow w_i^t \beta.$$

More general setting:

- Do not know there is one expert who makes no mistakes.
- Weighted Majority Algorithm:
    - at any time, expert i has weight w_i^t.
    - originally, $w_i^t = 1, \forall i \in [1, N]$.
    - prediction according to weighted majority.
    - weight of each wrong expert updated ($\beta \in (0, 1)$)

    $$w_i^{t+1} \leftarrow w_i^t \beta.$$

- Theorem: let m_i^t be the number of mistakes made by expert i till time t and m^t be the number of mistakes made by the WM algorithm. Then

    $$m_t \leq \frac{\log_2 N + m_i^t \log_2 \frac{1}{\beta}}{\log_2 \frac{2}{1+\beta}}.$$

    - $m_t \leq O(\log_2 N) + \text{constant} \times \text{best expert}.$

Remarks:

- Linear dependence on the best expert

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions
- Independent of the choice of sequence of outcomes

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions
- Independent of the choice of sequence of outcomes
- Mild dependence on the number of the experts

General Setting

- For $t = 1$ to T do
  - receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y$, $i \in [1, N]$
  - predict $\hat{y}_t \in Y$
  - receive label $y_t \in Y$
  - incur loss $L(y_t, \hat{y}_t)$

General Setting

- For $t = 1$ to T do
  - ▸ receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y$, $i \in [1, N]$
  - ▸ predict $\hat{y}_t \in Y$
  - ▸ receive label $y_t \in Y$
  - ▸ incur loss $L(y_t, \hat{y}_t)$

- Objective: minimize regret, i.e., difference of total loss incurred and that of best exeprt:

$$Regret(T) = \sum_{t=1}^{T} L(y_t, \hat{y}_t) - \min_{i \in [1,N]} \sum_{t=1}^{T} L(y_t, \hat{y}_{t,i}).$$

General Setting

- For $t = 1$ to T do
  - receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y$, $i \in [1, N]$
  - predict $\hat{y}_t \in Y$
  - receive label $y_t \in Y$
  - incur loss $L(y_t, \hat{y}_t)$

- Objective: minimize regret, i.e., difference of total loss incurred and that of best exeprt:

$$Regret(T) = \sum_{t=1}^{T} L(y_t, \hat{y}_t) - \min_{i \in [1,N]} \sum_{t=1}^{T} L(y_t, \hat{y}_{t,i}).$$

- Goal: obtain sublinear regret to the best expert. (top 10 experts, any competitor, etc.)

$$\lim_{T \to \infty} \frac{Regret(T)}{T} = 0.$$

An Optimal Bound

- Exponential weighted average
  - weight update:
    $$w_i^{t+1} \leftarrow w_i^t e^{-\eta_t L_{t,i}},$$
    where $L_{t,i}$ is total loss incurred by expert i till time t
  - prediction:
    $$\hat{y}_t = \frac{\sum_{i=1}^N w_i^t y_{t,i}}{\sum_{i=1}^N w_i^t}.$$

An Optimal Bound

- Exponential weighted average
  - weight update:
    $$w_i^{t+1} \leftarrow w_i^t e^{-\eta_t L_{t,i}},$$
    where $L_{t,i}$ is total loss incurred by expert i till time t
  - prediction:
    $$\hat{y}_t = \frac{\sum_{i=1}^N w_i^t y_{t,i}}{\sum_{i=1}^N w_i^t}.$$

- Theorem: Assume L is convex in its first argument and takes values in $[0,1]$. Then for any T

$$Regret(T) \leq \frac{\sqrt{2}}{\sqrt{2}-1}\sqrt{(T/2)\log_2 N} + \sqrt{\log_2 N/2}.$$

Remarks on the theorem

- When the loss is between $[0, 1]$

$$Regret(T) = O(\sqrt{(\log N)T}).$$

Remarks on the theorem

- When the loss is between $[0, 1]$

$$Regret(T) = O(\sqrt{(\log N)T}).$$

- The bound can be improved if more information is known, e.g., if the loss functions are special, L_2 loss, L_1 loss,

Remarks on the theorem

- When the loss is between $[0, 1]$

$$Regret(T) = O(\sqrt{(\log N)T}).$$

- The bound can be improved if more information is known, e.g., if the loss functions are special, L_2 loss, L_1 loss,
- In general, the bound is asymptotically optimal.

Active research area: literatures

- Littlestone and Warmuth develop weighted majority algorithm
- Cesa-Bianchi et al. obtain minimax regret bounded by $\sqrt{(T/2)\log N}$.
- Cesa-Bianchi and Lugosi connected the regret with the Rademacher Complexity of the set of experts.
- Even-Dar et al. considered regret to the average
- Rakhlin et al. developed Sequential Rademacher Complexity to study online learning.
- . . .

A naive prediction problem (T.Cover, 1965)

- Predict $y_1, y_2, \ldots,$ where $y_t \in \{-1, 1\}$ (stock goes up or down).

A naive prediction problem (T.Cover, 1965)

- Predict $y_1, y_2, \ldots$, where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell $\hat{p}_1$ shares of stocks) according to some algorithm $\hat{P}$.

A naive prediction problem (T.Cover, 1965)

- Predict $y_1, y_2, \ldots$, where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell $\hat{p}_1$ shares of stocks) according to some algorithm $\hat{P}$.
- At each time t, investor loss (or gain) $-\hat{p}_t y_t$. Cumulative loss (or gain)

$$L(\hat{P}, y_T) := -\sum_{t=1}^{T} \hat{p}_t y_t$$

A naive prediction problem (T.Cover, 1965)

- Predict $y_1, y_2, \ldots$, where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell $\hat{p}_1$ shares of stocks) according to some algorithm $\hat{P}$.
- At each time t, investor loss (or gain) $-\hat{p}_t y_t$. Cumulative loss (or gain)

$$L(\hat{P}, y_T) := -\sum_{t=1}^{T} \hat{p}_t y_t$$

- Two naive experts: one always predicts 1 (always goes up); the other one always predicts -1(always goes down).

Cover proved

$$Regret(T) = \Theta(\sqrt{T}).$$

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy
- Typical Goal: minimize the regret with respect to the best performing expert.

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy
- Typical Goal: minimize the regret with respect to the best performing expert.
- More general goal: minimize the worst case value of

  ϕ(regret with respect to expert "1", regret with respect to expert "-1")

  at time T. ("Typical goal" is $\phi(x_1, x_2) = \min\{x_1, x_2\}$.)

Two very simple experts

Essentially, the problem is an optimal control problem:

Two very simple experts

Essentially, the problem is an optimal control problem:

- state space:

$$(x_1, x_2) = (\text{regret w.r.t "1" expert, regret w.r.t "-1" expert}).$$

Two very simple experts

Essentially, the problem is an optimal control problem:

- state space:

  $$(x_1, x_2) = (\text{regret w.r.t ``1'' expert}, \text{regret w.r.t ``-1'' expert}).$$

- control: prediction $\hat{p}_t \in [-1, 1]$ (investor's purchase bounded by 1)

Two very simple experts

Essentially, the problem is an optimal control problem:

- state space:

$$(x_1, x_2) = (\text{regret w.r.t "1" expert, regret w.r.t "-1" expert}).$$

- control: prediction $\hat{p}_t \in [-1, 1]$ (investor's purchase bounded by 1)
- value function: $v(x, t) = $ optimal time T result, starting from relative regrets $x = (x_1, x_2)$ at time t.

Dynamical programming principle:

$$v(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b = \pm 1} v(\text{new position}, t + 1)$$

$$= \min_{|p| \leq 1} \max_{b = \pm 1} v(x_1 + b(1 - p), x_2 - b(1 + p), t + 1)$$

for $t < T$, with final-time condition $v(x, T) = \phi(x)$.

Dynamical Programming Principle

Recall: $(x_1, x_2) = ($regret w.r.t. "1" expert, regret w.r.t. "-1" expert$)$, where regret = (investor's loss) - (expert's loss).

If the investor buys p shares and the market goes up, investor gains p, the "+" expert gains 1, the "-1" expert gains -1. So the state moves from (x_1, x_2) to $(x_1 + (1 - p), x_2 + (-1 - p))$.

Similarly, if investor buys p shares and market goes down, state move from (x_1, x_2) to $(x_1 - (1 - p), x_2 - (-1 - p))$.

Hence the dynamical programming principle:

$$v(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b = \pm 1} v(x_1 + b(1 - p), x_2 - b(1 + p), t + 1).$$

- Want to know how the regret accumulates as $T \to \infty$.

- Want to know how the regret accumulates as $T \rightarrow \infty$.
- To access this question, it is natural to rescale the problem and look for a continuum limit.

- Want to know how the regret accumulates as $T \to \infty$.
- To access this question, it is natural to rescale the problem and look for a continuum limit.
- Our scaling is like the passage from random walk to Brownian motion.
  - Let $\xi_1, \xi_2, \ldots$ be i.i.d. random variable with mean 0 and variance 1. For each n, define a continuous-time stochastic processes $\{W_n(t)\}$

$$W_n(t) = \frac{1}{\sqrt{n}} \sum_{1 \leq j \leq [nt]} \xi_j.$$

Then $W_n(t) \to W_t$, where W_t is a standard Brownian motion.

- So we consider a scaled version of the problem: stock moves $\pm\epsilon$, time steps are ϵ^2, and the scaling of regret is ϵ^2. The value function is still the optimal time-T result. The dynamical programming principle becomes

$$w^\epsilon(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b = \pm 1} w^\epsilon(x_1 + \epsilon b(1-p), x_2 - \epsilon b(1+p), t + \epsilon^2)$$

- So we consider a scaled version of the problem: stock moves $\pm\epsilon$, time steps are ϵ^2, and the scaling of regret is ϵ^2. The value function is still the optimal time-T result. The dynamical programming principle becomes

$$w^\epsilon(x_1, x_2, t) = \min_{|p|\leq 1} \max_{b=\pm 1} w^\epsilon(x_1 + \epsilon b(1-p), x_2 - \epsilon b(1+p), t + \epsilon^2)$$

- We expect as $\epsilon \to 0$, $w^\epsilon(x, t) \to w(x, t)$ and w solves some PDEs.

The partial differential equation is the Hamilton-Jacobi-Bellman equation associate with the optimal control problem. Formal derivation

1. Use the Taylor's expansion to estimate

$$w(x_1 + \epsilon b(1 - p), x_2 - \epsilon(1 + p), t + \epsilon^2).$$

2. Investor choose p to make $O(\epsilon)$ terms vanishes; this gives

$$p = \frac{\partial_1 w - \partial_2 w}{\partial_1 w + \partial_2 w}.\text{(optimal strategy)}$$

3. The $O(\epsilon^2)$ are insensitive to $b = \pm 1$; they give the nonlinear PDE

$$w_t + 2 < D^2 w \frac{\nabla^\perp w}{\partial_1 w + \partial_2 w}, \frac{\nabla^\perp w}{\partial_1 w + \partial_2 w} >= 0,$$

where $\nabla^\perp w = (\partial_2 w, -\partial_1 w)$. Prescribe the final value $w = \phi$ at T.

Conclusion

- if ϕ smooth

$$\|w^\epsilon(x, t) - w(x, t)\| \leq C\epsilon,$$

  where C independent of t.
- if $\phi(x) = \min\{x_1, x_2\}$

$$\|w^\epsilon(x, t) - w(x, t)\| \leq C\epsilon \log \epsilon.$$

- Equation can be solved explicitly.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.
- One can consider more complex situation, e.g., more experts, more intelligent experts, etc.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.
- One can consider more complex situation, e.g., more experts, more intelligent experts, etc.
- Dynamical programming principle can be hard to derive.

- In the previous example, we derive the PDE from Hamilton-Jacobi-Bellman equation from stochastic control.

- In the previous example, we derive the PDE from Hamilton-Jacobi-Bellman equation from stochastic control.
- In the discrete case, Markov decision process (MDP) is a discrete time stochastic control process. (Reinforcement Learning). The value function satisfies a discrete Bellman equation.

- In the previous example, we derive the PDE from Hamilton-Jacobi-Bellman equation from stochastic control.
- In the discrete case, Markov decision process (MDP) is a discrete time stochastic control process. (Reinforcement Learning). The value function satisfies a discrete Bellman equation.
  - Given a policy π with infinite horizon for any state $x \in X$

  $$V_\pi(x) = R(x, \pi(x)) + \gamma \sum_{x' \in X} P(x'|x, \pi(x)) V_\pi(x')$$

  - Optimal policy

  $$V^*(x) = \max_{a \in A} \{R(x, a) + \gamma \sum_{x' \in X} P(x'|x, a) V_\pi(x')\}$$

- Continuous time and state: PDEs.
  - ▸ Dynamics of the state

  $$dX_s = b(X_s, a_s)\, ds + \sigma(X_s, a_s)\, dW_s, \tag{1}$$

  where W_s is Brownian motion and a_s is control process.

- Continuous time and state: PDEs.
  - ▶ Dynamics of the state

$$dX_s = b(X_s, a_s) \, ds + \sigma(X_s, a_s) \, dW_s, \tag{1}$$

  where W_s is Brownian motion and a_s is control process.
  - ▶ Reward

$$R(t, x, a) = \mathbb{E}\big[\int_t^T f(X_s^{t,x}, a_s) \, ds + g(X_T^{t,x})\big],$$

  where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.

- Continuous time and state: PDEs.
  - ▶ Dynamics of the state

$$dX_s = b(X_s, a_s)\, ds + \sigma(X_s, a_s)\, dW_s, \qquad (1)$$

  where W_s is Brownian motion and a_s is control process.
  - ▶ Reward

$$R(t, x, a) = \mathbb{E}\big[\int_t^T f(X_s^{t,x}, a_s)\, ds + g(X_T^{t,x})\big],$$

  where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.
  - ▶ Value function

$$v(t, x) = \sup_{a \in \mathcal{A}} R(t, x, a).$$

- Continuous time and state: PDEs.
  - ▶ Dynamics of the state

  $$dX_s = b(X_s, a_s)\, ds + \sigma(X_s, a_s)\, dW_s, \qquad (1)$$

  where W_s is Brownian motion and a_s is control process.
  - ▶ Reward

  $$R(t, x, a) = \mathbb{E}\big[\int_t^T f(X_s^{t,x}, a_s)\, ds + g(X_T^{t,x})\big],$$

  where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.
  - ▶ Value function

  $$v(t, x) = \sup_{a \in \mathcal{A}} R(t, x, a).$$

  - ▶ HJB equation

  $$\frac{\partial v}{\partial t} + \sup_{a \in \mathcal{A}} (L_a v + f(x, a)) = 0 \quad \text{in} [0, T] \times \mathbb{R}^d,$$

  where L_a is the infinitesimal generator of the process (1).

Second-order Linear Elliptic PDEs

A classical example is

$$\Delta u(x) = \sum_{i=1}^{d} \frac{\partial^2 u}{\partial x_i^2} = 0.$$

Let W_t be a Brownian motion starting from x. Then

$$\lim_{t \to 0} \frac{E^x u(W_t) - u(x)}{t} = \frac{1}{2}\Delta u.$$

General diffusion process

$$L_a u = \frac{1}{2} Tr(\sigma \sigma^T D^2 u) = \frac{1}{2} \sum_{i,j=1}^{d} a_{ij} D_{ij} u,$$

where $\{a_{ij}\}$ is some positive-definite matrix.

HJB equation

Fully nonlinear equation (v depends on (t, x))

$$\partial_t v + \sup_{a \in \mathcal{A}} (L_a v + f(x, a)) = 0,$$

where L_a is a second-order linear elliptic operator.

After deriving the HJB equation formally, the equation holds in a weak sense. The regularity of the solution is very low.

- Easy to find solution in the space of continuous function
- Equation holds in a very weak sense (viscosity solution).
- Important problem in the PDE community to prove the existence of classical solution.

Mathematical Result

Recall:
$$dX_s = b(X_s, a_s)\, ds + \sigma(X_s, a_s)\, dW_s.$$

Under some regularity assumption on the coefficients σ, b, then the solution to the HJB

$$\partial_t u + \sup_{a \in \mathcal{A}}(L_a u + f(x, a)) = 0,$$

is a classical solution. Rigorously, $u \in C^{2,\alpha}$ for some $\alpha > 0$.

The theoretical results helps to build numerical scheme to solve the HJB.

Jump processes and Nonlocal equations

Brownian motion is a continuous stochastic process. Consider more general process: Lévy processes (jump processes)

$$dX_s = b(X_s, a_s) \, ds + \sigma(X_s, a_s) \, dY_s,$$

where Y_s is a Lévy process. The infinitesimal generator of such processes is nonlocal

$$L_a u = \int_{\mathbb{R}^d} (u(x+y) + u(x-y) - 2u(x)) \frac{K_a(y)}{|y|^{d+\sigma}} \, dy,$$

and $\sigma \in (0, 2)$.

Theorem

Let v be a weak solution to

$$v_t + \sup_a (L_a v + f_a) = 0,$$

where L_a is a nonlocal operator defined previously. Then the solution v is a classical solution and derivatives of v are bounded.

Thank you.