

On PDE approach to some machine learning problems

Hong Zhang

Applied Math, Brown University


April 7, 2017

- 1 Introduction
- 2 Prediction with Expert Advice
 - A simple example
 - PDE approach of prediction problem
- 3 Stochastic Control
- 4 Elliptic PDEs

Partial differential equations (PDEs) are used to describe many phenomena. Many famous PDEs are from physics.

- Heat: heat equation, u temperature


$$\partial_t u = \sum_{i=1}^3 \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$



Partial differential equations (PDEs) are used to describe many phenomena. Many famous PDEs are from physics.

- Heat: heat equation, u temperature

$$\partial_t u = \sum_{i=1}^3 \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$



- Fluid mechanics: Navier-Stokes equation


$$\partial_t u + (u \cdot \nabla) u - \Delta u = -\nabla p.$$



Partial differential equations (PDEs) are used to describe many phenomena. Many famous PDEs are from physics.

- Heat: heat equation, u temperature

$$\partial_t u = \sum_{i=1}^3 \frac{\partial^2 u}{\partial x_i^2} = \Delta u$$



- Fluid mechanics: Navier-Stokes equation

$$\partial_t u + (u \cdot \nabla) u - \Delta u = -\nabla p.$$



- Quantum mechanics, electrodynamics, image processing, . . .

Probability and PDEs




Figure: Brownian Motion

W_t : a Brownian motion with $W_0 = x$. W hits boundary at W_τ .
 $u(x) = \mathbb{E}^x[g(W_\tau)]$. Then

$$\Delta u = 0,$$

and

$$u = g \quad \text{on} \quad S_1.$$

Online Learning: Examples



Figure: Spam Detection

- emails arrive one by one

Online Learning: Examples



Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: **spam** or **not**

Online Learning: Examples



Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: **spam** or **not**
- sometimes get **feedback** from users

Online Learning: Examples



Figure: Spam Detection

- emails arrive one by one
- spam detector makes prediction for each of them: **spam or not**
- sometimes get **feedback** from users
- make better detection next time

Online Learning: Example 2



- users visit Netflix /Amazon/...

Figure: Recommendation

Online Learning: Example 2



Figure: Recommendation

- users visit Netflix /Amazon/...
- movies/products/... are recommended

Online Learning: Example 2



Figure: Recommendation

- users visit Netflix /Amazon/...
- movies/products/... are recommended
- users **click or not**

Online Learning: Example 2



Figure: Recommendation

- users visit Netflix /Amazon/...
- movies/products/... are recommended
- users **click or not**
- make better recommendation next time

Why Online Learning?

- Batch learning:
 - ▶ distribution fixed over time (training and test)
 - ▶ i.i.d. assumptions
 - ▶ a fixed predictor

Why Online Learning?

- Batch learning:
 - ▶ distribution fixed over time (training and test)
 - ▶ i.i.d. assumptions
 - ▶ a fixed predictor
- Online learning:
 - ▶ no distribution assumption
 - ▶ suitable for changing or adversarial environment
 - ▶ a natural model for many applications, mixing training and test
 - ▶ memory efficient

Prediction With Expert Advices

- Need: to make a decision every day
 - ▶ which stock to invest in
 - ▶ which route to drive home
 - ▶ ...



Prediction With Expert Advices

- Need: to make a decision every day
 - ▶ which stock to invest in
 - ▶ which route to drive home
 - ▶ ...
- Given: advice from a set of experts
 - ▶ financial adviser
 - ▶ Route 1, Route I-95, Route 295,...
 - ▶ ...



Prediction With Expert Advices

- Need: to make a decision every day
 - ▶ which stock to invest in
 - ▶ which route to drive home
 - ▶ ...
- Given: advice from a set of experts
 - ▶ financial adviser
 - ▶ Route 1, Route I-95, Route 295,...
 - ▶ ...
- Question: How to combine advice and make smart choices?



Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.

Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.
 - ▶ N experts provide advice $(f_{1,t}, f_{2,t}, \dots, f_{N,t})$

Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.
 - ▶ N experts provide advice $(f_{1,t}, f_{2,t}, \dots, f_{N,t})$
 - ▶ A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t

Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.
 - ▶ N experts provide advice $(f_{1,t}, f_{2,t}, \dots, f_{N,t})$
 - ▶ A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
 - ▶ True y_t is revealed and find if $\hat{p}_t = y_t$.

Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.
 - ▶ N experts provide advice $(f_{1,t}, f_{2,t}, \dots, f_{N,t})$
 - ▶ A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
 - ▶ True y_t is revealed and find if $\hat{p}_t = y_t$.
 - ▶ Know that one of the experts makes no mistakes.

Gentle Start: a simple example

- Set up:
 - ▶ Predict an unknown sequence y_1, y_2, \dots , where $y_i \in \{0, 1\}$.
 - ▶ N experts provide advice $(f_{1,t}, f_{2,t}, \dots, f_{N,t})$
 - ▶ A forecaster makes her guess $\hat{p}_t \in \{0, 1\}$ for y_t
 - ▶ True y_t is revealed and find if $\hat{p}_t = y_t$.
 - ▶ Know that one of the experts makes no mistakes.
- Goal: bound the number of the mistakes made by the forecaster

Halving Algorithm

Algorithm:

- at time t , expert i has weight w_i^t

Halving Algorithm

Algorithm:

- at time t , expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$

Halving Algorithm

Algorithm:

- at time t , expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority

Halving Algorithm

Algorithm:

- at time t , expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority
- update each wrong expert $w_i^t = 0$

Halving Algorithm

Algorithm:

- at time t , expert i has weight w_i^t
- originally, $w_i^0 = 1, \forall i \in [1, N]$
- prediction according to the weighted majority
- update each wrong expert $w_i^t = 0$

Conclusion: the number of mistakes made by the forecaster is bounded by $\lfloor \log_2 N \rfloor$.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.
- When the forecaster make her m th mistakes, at least half of the experts have been always correct make their first mistake

$$w_m \leq \frac{w_{m-1}}{2} \implies w_m \leq \frac{w_0}{2^m}.$$

Proof

:

- let w_m be the sum of the weight of all experts after the forecaster has made m mistakes.
- Initially, $w_0 = N$.
- When the forecaster make her m th mistakes, at least half of the experts have been always correct make their first mistake

$$w_m \leq \frac{w_{m-1}}{2} \implies w_m \leq \frac{w_0}{2^m}.$$

- $w_i = 1$ for some i

$$m \leq \lfloor \log_2 N \rfloor.$$

More general setting:

- Do not know there is one expert who makes no mistakes.

More general setting:

- Do not know there is one expert who makes no mistakes.
- Weighted Majority Algorithm:
 - ▶ at any time, expert i has weight w_i^t .
 - ▶ originally, $w_i^t = 1, \forall i \in [1, N]$.
 - ▶ prediction according to weighted majority.
 - ▶ weight of each wrong expert updated ($\beta \in (0, 1)$)

$$w_i^{t+1} \leftarrow w_i^t \beta.$$

More general setting:

- Do not know there is one expert who makes no mistakes.
- Weighted Majority Algorithm:
 - ▶ at any time, expert i has weight w_i^t .
 - ▶ originally, $w_i^t = 1, \forall i \in [1, N]$.
 - ▶ prediction according to weighted majority.
 - ▶ weight of each wrong expert updated ($\beta \in (0, 1)$)

$$w_i^{t+1} \leftarrow w_i^t \beta.$$

- Theorem: let m_i^t be the number of mistakes made by expert i till time t and m^t be the number of mistakes made by the WM algorithm. Then

$$m_t \leq \frac{\log_2 N + m_i^t \log_2 \frac{1}{\beta}}{\log_2 \frac{2}{1+\beta}}.$$

- ▶ $m_t \leq O(\log_2 N) + \text{constant} \times \text{best expert}.$

Remarks:

- Linear dependence on the best expert

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions
- Independent of the choice of sequence of outcomes

Remarks:

- Linear dependence on the best expert
- Independent of the number of predictions
- Independent of the choice of sequence of outcomes
- Mild dependence on the number of the experts

General Setting

- For $t = 1$ to T do
 - ▶ receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y, i \in [1, N]$
 - ▶ predict $\hat{y}_t \in Y$
 - ▶ receive label $y_t \in Y$
 - ▶ incur loss $L(y_t, \hat{y}_t)$

General Setting

- For $t = 1$ to T do
 - ▶ receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y, i \in [1, N]$
 - ▶ predict $\hat{y}_t \in Y$
 - ▶ receive label $y_t \in Y$
 - ▶ incur loss $L(y_t, \hat{y}_t)$
- **Objective:** minimize regret, i.e., difference of total loss incurred and that of best expeprt:

$$\text{Regret}(T) = \sum_{t=1}^T L(y_t, \hat{y}_t) - \min_{i \in [1, N]} \sum_{t=1}^T L(y_t, \hat{y}_{t,i}).$$

General Setting

- For $t = 1$ to T do
 - ▶ receive instance $x_t \in X$ and advice $\hat{y}_{t,i} \in Y, i \in [1, N]$
 - ▶ predict $\hat{y}_t \in Y$
 - ▶ receive label $y_t \in Y$
 - ▶ incur loss $L(y_t, \hat{y}_t)$
- **Objective:** minimize regret, i.e., difference of total loss incurred and that of best expert:

$$\text{Regret}(T) = \sum_{t=1}^T L(y_t, \hat{y}_t) - \min_{i \in [1, N]} \sum_{t=1}^T L(y_t, \hat{y}_{t,i}).$$

- Goal: obtain **sublinear regret** to the **best expert**. (top 10 experts, any competitor, etc.)

$$\lim_{T \rightarrow \infty} \frac{\text{Regret}(T)}{T} = 0.$$

An Optimal Bound

- Exponential weighted average

- ▶ weight update:

$$w_i^{t+1} \leftarrow w_i^t e^{-\eta_t L_{t,i}},$$

where $L_{t,i}$ is total loss incurred by expert i till time t

- ▶ prediction:

$$\hat{y}_t = \frac{\sum_{i=1}^N w_i^t y_{t,i}}{\sum_{i=1}^N w_i^t}.$$

An Optimal Bound

- Exponential weighted average

- ▶ weight update:

$$w_i^{t+1} \leftarrow w_i^t e^{-\eta L_{t,i}},$$

where $L_{t,i}$ is total loss incurred by expert i till time t

- ▶ prediction:

$$\hat{y}_t = \frac{\sum_{i=1}^N w_i^t y_{t,i}}{\sum_{i=1}^N w_i^t}.$$

- Theorem: Assume L is convex in its first argument and takes values in $[0, 1]$. Then for any T

$$\text{Regret}(T) \leq \frac{\sqrt{2}}{\sqrt{2}-1} \sqrt{(T/2) \log_2 N} + \sqrt{\log_2 N/2}.$$

Remarks on the theorem

- When the loss is between $[0, 1]$

$$\text{Regret}(T) = O(\sqrt{(\log N) T}).$$

Remarks on the theorem

- When the loss is between $[0, 1]$

$$\text{Regret}(T) = O(\sqrt{(\log N) T}).$$

- The bound can be improved if more information is known, e.g., if the loss functions are special, L_2 loss, L_1 loss,

Remarks on the theorem

- When the loss is between $[0, 1]$

$$\text{Regret}(T) = O(\sqrt{(\log N) T}).$$

- The bound can be improved if more information is known, e.g., if the loss functions are special, L_2 loss, L_1 loss,
- In general, the bound is asymptotically optimal.

Active research area: literatures

- Littlestone and Warmuth develop [weighted majority algorithm](#)
- Cesa-Bianchi et al. obtain minimax regret bounded by $\sqrt{(T/2) \log N}$.
- Cesa-Bianchi and Lugosi connected the regret with the Rademacher Complexity of the set of experts.
- Even-Dar et al. considered regret to the average
- Rakhlin et al. developed [Sequential Rademacher Complexity](#) to study online learning.
- ...

A naive prediction problem (T.Cover, 1965)

- Predict y_1, y_2, \dots , where $y_t \in \{-1, 1\}$ (stock goes up or down).

A naive prediction problem (T.Cover, 1965)

- Predict y_1, y_2, \dots , where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell \hat{p}_1 shares of stocks) according to some algorithm \hat{P} .

A naive prediction problem (T.Cover, 1965)

- Predict y_1, y_2, \dots , where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell \hat{p}_1 shares of stocks) according to some algorithm \hat{P} .
- At each time t , investor loss (or gain) $-\hat{p}_t y_t$. Cumulative loss (or gain)

$$L(\hat{P}, y_T) := - \sum_{t=1}^T \hat{p}_t y_t$$

A naive prediction problem (T.Cover, 1965)

- Predict y_1, y_2, \dots , where $y_t \in \{-1, 1\}$ (stock goes up or down).
- Forecaster makes prediction $\hat{p}_t \in [-1, 1]$ (buy/sell \hat{p}_1 shares of stocks) according to some algorithm \hat{P} .
- At each time t , investor loss (or gain) $-\hat{p}_t y_t$. Cumulative loss (or gain)

$$L(\hat{P}, y_T) := - \sum_{t=1}^T \hat{p}_t y_t$$

- Two naive experts: one always predicts 1 (always goes up); the other one always predicts -1 (always goes down).

Cover proved

$$\text{Regret}(T) = \Theta(\sqrt{T}).$$

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy
- Typical Goal: minimize the worst-case regret with respect to the best performing expert.

PDE approach

This problem was considered by Kohn using PDE technique.

- Instead of proposing an algorithm and then proving the asymptotic bound, PDE method focuses on the optimal strategy
- Typical Goal: minimize the worst-case regret with respect to the best performing expert.
- More general goal: minimize the worst case value of

ϕ (regret with respect to expert “1”, regret with respect to expert “-1”)

at time T . (“Typical goal” is $\phi(x_1, x_2) = \max\{x_1, x_2\}$.)

Two very simple experts

Essentially, the problem is an optimal control problem:

Two very simple experts

Essentially, the problem is an optimal control problem:

- state space:

$$(x_1, x_2) = (\text{regret w.r.t. expert "1"}, \text{regret w.r.t. expert "-1"}).$$

Two very simple experts

Essentially, the problem is an optimal control problem:

- **state space:**

$$(x_1, x_2) = (\text{regret w.r.t. expert "1", regret w.r.t. expert "-1"}).$$

- **control:** prediction $\hat{p}_t \in [-1, 1]$ (investor's purchase bounded by 1)

Two very simple experts

Essentially, the problem is an optimal control problem:

- **state space:**

$$(x_1, x_2) = (\text{regret w.r.t. expert "1", regret w.r.t. expert "-1"}).$$

- **control:** prediction $\hat{p}_t \in [-1, 1]$ (investor's purchase bounded by 1)
- **value function:** $v(x, t)$ = optimal time T result, starting from relative regrets $x = (x_1, x_2)$ at time t .

Dynamical programming principle:

$$\begin{aligned} v(x_1, x_2, t) &= \min_{|p| \leq 1} \max_{b=\pm 1} v(\text{new position}, t+1) \\ &= \min_{|p| \leq 1} \max_{b=\pm 1} v(x_1 + b(1-p), x_2 - b(1+p), t+1) \end{aligned}$$

for $t < T$, with final-time condition $v(x, T) = \phi(x)$.

Dynamical Programming Principle

Recall: $(x_1, x_2) = (\text{regret w.r.t. "1" expert, regret w.r.t. "-1" expert})$,
where $\text{regret} = (\text{investor's loss}) - (\text{expert's loss})$.

If the investor buys p shares and the market goes up, investor loss $-p$, the "+" expert loss -1 , the "-1" expert loss 1. So the state moves from (x_1, x_2) to $(x_1 - (1 - p), x_2 + (1 + p))$.

Similarly, if investor buys p shares and market goes down, state move from (x_1, x_2) to $(x_1 + (1 - p), x_2 - (1 + p))$.

Hence the dynamical programming principle:

$$v(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b=\pm 1} v(x_1 + b(1 - p), x_2 - b(1 + p), t + 1).$$

- Want to know how the regret accumulates as $T \rightarrow \infty$.

- Want to know how the regret accumulates as $T \rightarrow \infty$.
- To access this question, it is natural to rescale the problem and look for a continuum limit.

- Want to know how the regret accumulates as $T \rightarrow \infty$.
- To access this question, it is natural to rescale the problem and look for a continuum limit.
- Our scaling is like the passage from random walk to Brownian motion.
 - ▶ Let ξ_1, ξ_2, \dots be i.i.d. random variable with mean 0 and variance 1. For each n , define a continuous-time stochastic processes $\{W_n(t)\}$

$$W_n(t) = \frac{1}{\sqrt{n}} \sum_{1 \leq j \leq [nt]} \xi_j.$$

Then $W_n(t) \rightarrow W_t$, where W_t is a standard Brownian motion.

- So we consider a **scaled** version of the problem: stock moves $\pm\epsilon$, time steps are ϵ^2 , and the scaling of regret is ϵ^2 . The value function is still the optimal time-T result. The dynamical programming principle becomes

$$w^\epsilon(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b=\pm 1} w^\epsilon(x_1 + \epsilon b(1-p), x_2 - \epsilon b(1+p), t + \epsilon^2)$$

- So we consider a **scaled** version of the problem: stock moves $\pm\epsilon$, time steps are ϵ^2 , and the scaling of regret is ϵ^2 . The value function is still the optimal time-T result. The dynamical programming principle becomes

$$w^\epsilon(x_1, x_2, t) = \min_{|p| \leq 1} \max_{b=\pm 1} w^\epsilon(x_1 + \epsilon b(1-p), x_2 - \epsilon b(1+p), t + \epsilon^2)$$

- We expect as $\epsilon \rightarrow 0$, $w^\epsilon(x, t) \rightarrow w(x, t)$ and w solves some PDEs.

The partial differential equation is the **Hamilton-Jacobi-Bellman** equation associate with the optimal control problem. Formal derivation

- 1 Use the Taylor's expansion to estimate

$$w(x_1 + \epsilon b(1 - p), x_2 - \epsilon(1 + p), t + \epsilon^2).$$

- 2 Investor choose p to make $O(\epsilon)$ terms vanishes; this gives

$$p = \frac{\partial_1 w - \partial_2 w}{\partial_1 w + \partial_2 w} \cdot (\text{optimal strategy})$$

- 3 The $O(\epsilon^2)$ are insensitive to $b = \pm 1$; they give the nonlinear PDE

$$w_t + 2 < D^2 w \frac{\nabla^\perp w}{\partial_1 w + \partial_2 w}, \frac{\nabla^\perp w}{\partial_1 w + \partial_2 w} > = 0,$$

where $\nabla^\perp w = (\partial_2 w, -\partial_1 w)$. Prescribe the final value $w = \phi$ at T .

Conclusion

- if ϕ smooth

$$\|w^\epsilon(x, t) - w(x, t)\| \leq C\epsilon,$$

where C independent of t .

- if $\phi(x) = \max\{x_1, x_2\}$

$$\|w^\epsilon(x, t) - w(x, t)\| \leq C\epsilon \log \epsilon.$$

- Equation can be solved explicitly.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.


Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.
- One can consider more complex situation, e.g., more experts, more intelligent experts, etc.

Discussion

- Stock prediction problem has a continuous-time limit. The PDE has a rather explicit solution.
- Problem is considered as a deterministic two-person game.
- One can consider more complex situation, e.g., more experts, more intelligent experts, etc.
- Dynamical programming principle can be hard to derive.

Reinforcement Learning



S: Set of states

A: Set of actions

R: Reward function. $\forall s_t, a_t, r_t = R(s_t, a_t) \in \mathbb{R}$

- In the previous example, we derive the PDE from **Hamilton-Jacobi-Bellman** equation of an optimal control problem.

- In the previous example, we derive the PDE from **Hamilton-Jacobi-Bellman** equation of an optimal control problem.
- In the discrete case, Markov decision process (MDP) is a discrete time stochastic control process. (Reinforcement Learning). The value function $v : S \rightarrow \mathbb{R}$ satisfies a discrete Bellman equation.

- In the previous example, we derive the PDE from **Hamilton-Jacobi-Bellman** equation of a optimal control problem.
- In the discrete case, Markov decision process (MDP) is a discrete time stochastic control process. (Reinforcement Learning). The value function $v : S \rightarrow \mathbb{R}$ satisfies a discrete Bellman equation.
 - ▶ Given a policy π with infinite horizon for any state $x \in X$

$$V_{\pi}(x) = R(x, \pi(x)) + \gamma \sum_{x' \in X} P(x'|x, \pi(x)) V_{\pi}(x')$$

- ▶ Optimal policy

$$V^*(x) = \max_{a \in A} \{ R(x, a) + \gamma \sum_{x' \in X} P(x'|x, a) V_{\pi}(x') \}$$

- Continuous time and state: PDEs.

- ▶ Dynamics of the state

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dW_s, \quad (1)$$

where W_s is Brownian motion and a_s is control process.

- Continuous time and state: PDEs.

- ▶ Dynamics of the state

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dW_s, \quad (1)$$

where W_s is Brownian motion and a_s is control process.

- ▶ Reward

$$R(t, x, a) = \mathbb{E} \left[\int_t^T f(X_s^{t,x}, a_s) ds + g(X_T^{t,x}) \right],$$

where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.

- Continuous time and state: PDEs.

- ▶ Dynamics of the state

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dW_s, \quad (1)$$

where W_s is Brownian motion and a_s is control process.

- ▶ Reward

$$R(t, x, a) = \mathbb{E} \left[\int_t^T f(X_s^{t,x}, a_s) ds + g(X_T^{t,x}) \right],$$

where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.

- ▶ Value function

$$v(t, x) = \sup_{a \in \mathcal{A}} R(t, x, a).$$

- Continuous time and state: PDEs.

- ▶ Dynamics of the state

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dW_s, \quad (1)$$

where W_s is Brownian motion and a_s is control process.

- ▶ Reward

$$R(t, x, a) = \mathbb{E} \left[\int_t^T f(X_s^{t,x}, a_s) ds + g(X_T^{t,x}) \right],$$

where $X^{t,x}$ is the solution of (1) starting from $X_t = x$.

- ▶ Value function

$$v(t, x) = \sup_{a \in \mathcal{A}} R(t, x, a).$$

- ▶ HJB equation

$$\frac{\partial v}{\partial t} + \sup_{a \in \mathcal{A}} (L_a v + f(x, a)) = 0 \quad \text{in } [0, T] \times \mathbb{R}^d,$$

where L_a is the infinitesimal generator of the process (1).

Second-order Linear Elliptic PDEs

A classical example is

$$\Delta u(x) = \sum_{i=1}^d \frac{\partial^2 u}{\partial x_i^2} = 0.$$

Let W_t be a Brownian motion starting from x . Then

$$\lim_{t \rightarrow 0} \frac{E^x u(W_t) - u(x)}{t} = \frac{1}{2} \Delta u.$$

General diffusion process

$$L_a u = \frac{1}{2} \text{Tr}(\sigma \sigma^T D^2 u) = \frac{1}{2} \sum_{i,j=1}^d a_{ij} D_{ij} u,$$

where $\{a_{ij}\}$ is some positive-definite matrix.

HJB equation

Fully nonlinear equation (v depends on (t, x))

$$\partial_t v + \sup_{a \in \mathcal{A}} (L_a v + f(x, a)) = 0,$$

where L_a is a second-order linear elliptic operator.

After deriving the HJB equation formally, the equation holds in a weak sense. The regularity of the solution is very low.

- Easy to find solution in the space of continuous function
- Equation holds in a very weak sense (viscosity solution).
- Important problem in the PDE community to prove the existence of classical solution.

Mathematical Result

Recall:

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dW_s.$$

Under some regularity assumption on the coefficients σ, b , then the solution to the HJB

$$\partial_t u + \sup_{a \in \mathcal{A}} (L_a u + f(x, a)) = 0,$$

is a classical solution. Rigorously, $u \in C^{2,\alpha}$ for some $\alpha > 0$.

The theoretical results helps to build numerical scheme to solve the HJB.

Jump processes and Nonlocal equations

Brownian motion is a continuous stochastic process. Consider more general process: Lévy processes (jump processes)

$$dX_s = b(X_s, a_s) ds + \sigma(X_s, a_s) dY_s,$$

where Y_s is a Lévy process. The infinitesimal generator of such processes is nonlocal

$$L_a u = \int_{\mathbb{R}^d} (u(x+y) + u(x-y) - 2u(x)) \frac{K_a(y)}{|y|^{d+\sigma}} dy,$$

and $\sigma \in (0, 2)$.

Theorem

Let v be a weak solution to

$$v_t + \sup_a (L_a v + f_a) = 0,$$

where L_a is a nonlocal operator defined previously. Then the solution v is a classical solution and derivatives of v are bounded.

Thank you.